

Problem Solving by Searching

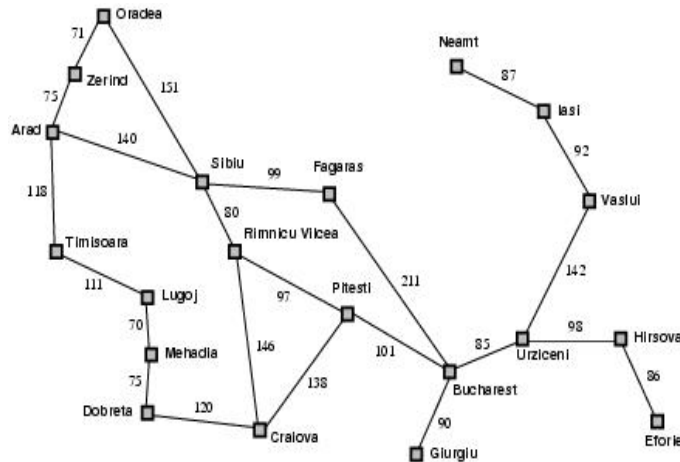


Outline



- Example problems
- Problem formulation
 - Example problems
- Searching for solutions
 - Tree search
 - Graph search
- Basic search algorithms
 - Uninformed

Example: Getting around in Romania



3

Example: Getting around in Romania



- On holiday in Romania; currently in Arad
 - Flight leaves tomorrow from Bucharest
- Formulate goal
 - Be in Bucharest
- Formulate problem
 - States: various cities
 - Actions: drive between cities
- Find solution
 - Sequence of cities; e.g. Arad, Sibiu, Fagaras, Bucharest, ...

4

Example: 8-puzzle



7	2	4
5		6
8	3	1


Start State

	1	2
3	4	5
6	7	8

Goal State

5

Example: 8-queens problem



- Constraint satisfaction problems (CSP)

6

Example Problems



- Toy problems
 - 8-puzzle
 - 8-queens
 - Rubik's cube
- Real-world problems
 - Route-finding problem: air travel planning, military operations planning, routing in computer networks
 - Touring problem, traveling salesperson problem
 - VLSI layout
 - Internet searching
 - Class scheduling

7

Problem types



- Goal-based agents
- Assumption on environment:
 - Observable: the agent knows current state
 - Static
 - Discrete: finitely many actions to choose
 - Deterministic
 - Agent knows the effects of its actions
- solution is a sequence of actions

8

Problem-solving agent



- Four general steps in problem solving:
 - Goal formulation
 - What are the successful world states
 - Problem formulation
 - What actions and states to consider given the goal
 - Design a representation that captures relevant aspects of the world
 - Search
 - Determine the possible sequence of actions that lead to the goal states and then choose the best sequence.
 - Execute
 - Given the solution perform the actions “eyes closed”

9

Problem formulation



- A search problem is defined by five components:
 - An **initial state**, e.g. “at Arad”
 - **Actions**: $ACTIONS(s)$ applicable actions in s
 - **Transition model** (or **Successor function**, or **Operators**)
 $RESULT(s,a)$: state that results from action a in state s
initial state + actions + transition model define **state space**, a directed **graph** in which the nodes are states and the arcs between nodes are actions
 - **Goal test**, can be
 - Explicit set of possible goal states, e.g. $x = 'at bucharest'$
 - Implicit, e.g. $checkmate(x)$, $NoDirt(x)$
 - **Path cost** (additive): assigns a numeric cost to each path
 - e.g. sum of distances, number of actions executed, ...
 - $c(s,a,s')$ is the **step cost**, assumed to be ≥ 0

A **solution** is a sequence of actions from initial to goal state.
Optimal solution has the lowest path cost.

10


Selecting a state space



- Real world is absurdly complex.
 - State space must be *abstracted* for problem solving.
- Good representations
 - preserve the relevant aspects of the problem
 - expose the relevant problem structure
 - abstracts away unimportant details

11

Example: vacuum world




Clean both squares with minimum number of actions

- States??
- Initial state??
- Actions?? Transition model??
- Goal test??
- Path cost??

12


Example: vacuum world



- States?? two locations, with or without dirt: $2 \times 2^2=8$ states.
- Initial state?? Any state can be initial
- Actions?? $\{Left, Right, Suck\}$
- Goal test?? Check whether squares are clean.
- Path cost?? 1 per action

13

Example: 8-puzzle




Start State

Goal State

- States??
- Initial state??
- Actions??
- Goal test??
- Path cost??

14

Example: 8-puzzle




- States?? Integer location of each tile, 10^5 states
- Initial state?? Any state can be initial
- Actions?? $\{Left, Right, Up, Down\}$
- Goal test?? Check whether goal configuration is reached
- Path cost?? 1 per move

[Note: optimal solution of n-Puzzle family is NP-complete]

15


Example: 8-queens problem



- States??
- Initial state??
- Actions??
- Goal test??
- Path cost??

16

Example: 8-queens problem




Incremental formulation

- States?? Any arrangement of 0 to 8 queens on the board
- Initial state?? No queens
- Actions?? Add queen in empty square
- Goal test?? 8 queens on board and none attacked
- Path cost?? None
10¹⁴ possible sequences to investigate

17

Example: 8-queens problem




Incremental formulation (alternative)

- States?? n ($0 \leq n \leq 8$) queens on the board, one per column in the n leftmost columns with no queen attacking another.
- Actions?? Add queen in leftmost empty column such that it is not attacking other queens
2057 possible sequences to investigate; 10⁵² for 100-queens

18

Example: 8-queens problem



complete-state formulation

- States?? 8 queens on the board, one per column
- Initial state?? any
- Actions?? Move a queen to a square in the same column
- Goal test?? 8 queens on board and none attacked
- Path cost?? None

19

Tree search algorithms



- How do we find the solutions of previous problems?
 - Search the state space (exhaustive search)
 - Generate a **search tree**
 - ROOT= initial state.
 - Branches are actions; Nodes and leafs generated through transition model
 - **Expanding** the current state: applying each legal action to the current state **generating** a new set of states

20

Tree search algorithms




```

function TREE-SEARCH(problem, strategy) return a solution or failure
  Initialize search tree to the initial state of the problem
  loop do
    if no candidates for expansion then return failure
    choose leaf node for expansion according to strategy
    if node contains goal state then return solution
    else expand the node and add resulting nodes to the search tree
  enddo
  
```

21

Simple tree search example



```


function TREE-SEARCH(problem, strategy) return a solution or failure
  Initialize search tree to the initial state of the problem
  do
    if no candidates for expansion then return failure
    choose leaf node for expansion according to strategy
    if node contains goal state then return solution
    else expand the node and add resulting nodes to the search tree
  enddo
  
```

22

Simple tree search example



(b) After expanding Arad



```


function TREE-SEARCH(problem, strategy) return a solution or failure
  Initialize search tree to the initial state of the problem
  do
    if no candidates for expansion then return failure
    choose leaf node for expansion according to strategy
    if node contains goal state then return solution
    else expand the node and add resulting nodes to the search tree
  enddo
  
```

23

Simple tree search example



(c) After expanding Sibiu




```

function TREE-SEARCH(problem, strategy) return a solution or failure
  Initialize search tree to the initial state of the problem
  do
    if no candidates for expansion then return failure
    choose leaf node for expansion according to strategy ← Determines search
    if node contains goal state then return solution process!!
    else expand the node and add resulting nodes to the search tree
  enddo
  
```

24

Implementation: states vs. nodes



- A *state* is a (representation of) a physical configuration
- A *node* is a bookkeeping data structure used to represent the search tree
 - $node = \langle state, parent-node, action, path-cost, \dots \rangle$
- **Fringe** = contains generated nodes which are not yet expanded = leaf nodes (aka, **Frontier**, **open list**)
 - Implemented as a queue

25

Tree search algorithm



```


function TREE-SEARCH(problem) return a solution or failure
    fringe ← INSERT(MAKE-NODE(problem.INITIAL-STATE))
    loop do
        if EMPTY?(fringe) then return failure
        node ← POP(fringe)
        if problem.GOAL-TEST(node.STATE)
            then return SOLUTION(node)
        for each action in problem.ACTIONS(node.STATE) do
            child ← Child-Node(problem, node, action)
            fringe ← INSERT(child, fringe)
    
```

Fringe implemented as a queue: FIFO, LIFO, priority queue

26

Repeated states

- Failure to detect repeated states can turn a solvable problem into unsolvable ones.



27

Graph search algorithm

- Closed list** (or **explored set**) stores all expanded nodes

function GRAPH-SEARCH(*problem*) **return** a solution or failure

closed ← an empty set

fringe ← INSERT(MAKE-NODE(*problem*.INITIAL-STATE))

loop do

if EMPTY?(*fringe*) **then return** failure

node ← POP(*fringe*)

if *problem*.GOAL-TEST(*node*.STATE)
 then return SOLUTION(*node*)

 add *node*.STATE to *closed*

for each action in *problem*.ACTIONS(*node*.STATE) **do**

child ← Child-Node(*problem*, *node*, action)

if *child*.STATE is not in *closed* or *fringe* **then**

fringe ← INSERT(*child*, *fringe*)

28

Uninformed search strategies



- (a.k.a. blind search) = use only information available in problem definition.
 - When strategies know whether one non-goal state is better than another → informed search.
- A strategy is defined by picking the **order of node expansion**:
 - Breadth-first search
 - Uniform-cost search
 - Depth-first search
 - Depth-limited search
 - Iterative deepening search.
 - Bidirectional search

29

Performance Measure




- Algorithm performance is measured in four ways:
 - **Completeness**; *Does it always find a solution if one exists?*
 - **Optimality**; *Does it always find the least-cost solution?*
 - **Time Complexity**; *Number of nodes generated?*
 - **Space Complexity**; *Number of nodes stored in memory during search?*
- Time and space complexity are measured in terms of problem difficulty defined by:
 - b – (maximum) **branching factor** of the search tree
 - d - depth of the shallowest solution
 - m - maximum length of the state space (may be ∞)

30

BF-search, an example



- Expand *shallowest* unexpanded node
- Implementation: *fringe* is a FIFO queue




31

BF-search, an example




- Expand *shallowest* unexpanded node
- Implementation: *fringe* is a FIFO queue



32

BF-search, an example


- Expand *shallowest* unexpanded node
- Implementation: *fringe* is a FIFO queue



33

BF-search, an example

- Expand *shallowest* unexpanded node
- Implementation: *fringe* is a FIFO queue



34

BF Graph search algorithm



```
function BF-SEARCH(problem) return a solution or failure
  closed ← an empty set
  fringe ← a FIFO queue with MAKE-NODE(problem.INITIAL-STATE)
  loop do
    if EMPTY?(fringe) then return failure
    node ← POP(fringe)
    add node.STATE to closed
    for each action in problem.ACTIONS(node.STATE) do
      child ← Child-Node(problem, node, action)
      if child.STATE is not in closed or fringe then
        if problem.GOAL-TEST(child.STATE)
          then return SOLUTION(child)
        fringe ← INSERT(child, fringe)
```

35

BF-search; evaluation



- Completeness:
 - *Does it always find a solution if one exists?*

36

BF-search; evaluation



- Completeness:

- *Does it always find a solution if one exists?*
- YES
 - Condition: If b is finite
 - (maximum number Of successor nodes is finite)

37

BF-search; evaluation



- Completeness:

- YES (if b is finite)

- Time complexity:

- Assume a state space where every state has b successors.
- Assume solution is at depth d

38

BF-search; evaluation



- Completeness:
 - YES (if b is finite)
- Time complexity (worst case tree search):
 - Assume a state space where every state has b successors.
 - root has b successors, each node at the next level has again b successors (total b^2), ...
 - Assume solution is at depth d
 - Total number of nodes generated in the worst case:

$$b + b^2 + b^3 + \dots + b^d = O(b^d)$$

39

BF-search; evaluation



- Completeness:
 - YES (if b is finite)
- Time complexity:
 - Total numb. of nodes generated:
$$b + b^2 + b^3 + \dots + b^d = O(b^d)$$
 - Graph search: more efficient, worst case proportional to the size of the state space (may be much smaller than $O(b^d)$).
- Space complexity?

40

BF-search; evaluation



- Completeness:
 - YES (if b is finite)
- Time complexity:
 - Total numb. of nodes generated:
$$b + b^2 + b^3 + \dots + b^d = O(b^d)$$
 - Graph search: proportional to the size of the state space (may be much smaller than $O(b^d)$).
- Space complexity: $O(b^d)$
 - Same as Time as each node is retained in memory
 - Tree search not save much space, use more time

41

BF-search; evaluation



- Completeness:
 - YES (if b is finite)
- Time complexity:
 - Total numb. of nodes generated: $O(b^d)$
 - Graph search: proportional to the size of the state space (may be much smaller than $O(b^d)$).
- Space complexity: $O(b^d)$
 - Same as Time as each node is retained in memory
- Optimality?
 - *Does it always find the least-cost solution?*

42

BF-search; evaluation



- Completeness:
 - YES (if b is finite)
- Time complexity:
 - Total numb. of nodes generated: $O(b^d)$
 - Graph search: more efficient, proportional to the size of the state space
- Space complexity: $O(b^d)$
 - Same as Time as each node is retained in memory
- Optimality?
 - *Does it always find the least-cost solution?*
 - Yes if unit cost per step.
 - No in general: actions have different cost.

43

BF-search; evaluation



- Two lessons:
 - Memory requirements are a bigger problem than its execution time.
 - Time: Exponential complexity search problems cannot be solved by uninformed search methods for any but the smallest instances.

DEPTH	NODES	TIME	MEMORY
2	110	0.11 milliseconds	107 kilobyte
4	11110	11 milliseconds	10.6 megabytes
6	10^6	1.1 seconds	1 gigabytes
8	10^8	2 minutes	103 gigabyte
10	10^{10}	3 hours	10 terabytes
12	10^{12}	13 days	1 petabytes
14	10^{14}	3.5 years	99 petabytes

$b=10$; 1 million nodes/second, 1000 bytes/node

44

Uniform-cost search



- Extension of BF-search:
 - Expand node with *lowest path cost* $g(n)$
- Implementation: *fringe* = priority queue ordered by path cost.
- UC-search is similar to BF-search when all step-costs are equal, but more costly

45

Uniform-cost Graph search



```
function Uniform-Cost-SEARCH(problem) return a solution or failure
  closed ← an empty set
  fringe ← a priority queue with MAKE-NODE(problem.INITIAL-STATE)
  loop do
    if EMPTY?(fringe) then return failure
    node ← POP(fringe)
    if problem.GOAL-TEST(node.STATE)
      then return SOLUTION(node)
    add node.STATE to closed
    for each action in problem.ACTIONS(node.STATE) do
      child ← Child-Node(problem, node, action)
      if child.STATE is not in closed or fringe then
        fringe ← INSERT(child, fringe)
      else if child.STATE is in fringe with higher path cost then
        replace that fringe node with child
```

46

Uniform-cost search




- Completeness:
 - YES, if step-cost $\geq \epsilon$ (small positive constant)
- Time complexity:
 - Expand nodes with cost less than C^* , the cost of the optimal solution.
 - Worst-case tree-search: $O(b^{1+C^*/\epsilon})$
- Space complexity:
 - Same as time complexity
- Optimality:
 - nodes expanded in order of increasing path cost.
 - YES, if complete.

47

DF-search, an example




- Expand *deepest* unexpanded node
- Implementation: *fringe* is a LIFO queue (=stack)



48

DF-search, an example


- Expand *deepest* unexpanded node
- Implementation: *fringe* is a LIFO queue (=stack)



49

DF-search, an example


- Expand *deepest* unexpanded node
- Implementation: *fringe* is a LIFO queue (=stack)



50

DF-search, an example


- Expand *deepest* unexpanded node
- Implementation: *fringe* is a LIFO queue (=stack)



51

DF-search, an example


- Expand *deepest* unexpanded node
- Implementation: *fringe* is a LIFO queue (=stack)



52

DF-search, an example


- Expand *deepest* unexpanded node
- Implementation: *fringe* is a LIFO queue (=stack)



53

DF-search, an example


- Expand *deepest* unexpanded node
- Implementation: *fringe* is a LIFO queue (=stack)



54

DF-search, an example


- Expand *deepest* unexpanded node
- Implementation: *fringe* is a LIFO queue (=stack)



55

DF-search, an example


- Expand *deepest* unexpanded node
- Implementation: *fringe* is a LIFO queue (=stack)



56

DF-search, an example


- Expand *deepest* unexpanded node
- Implementation: *fringe* is a LIFO queue (=stack)



57

DF-search, an example


- Expand *deepest* unexpanded node
- Implementation: *fringe* is a LIFO queue (=stack)



58

DF-search, an example

- Expand *deepest* unexpanded node
- Implementation: *fringe* is a LIFO queue (=stack)



59

DF-search; evaluation

- Completeness
 - Does it always find a solution if one exists?

60

DF-search; evaluation



- Completeness

- *Does it always find a solution if one exists?*
- NO: fails in infinite state spaces
 - Graph search complete in finite state spaces
 - Tree search complete in finite depth trees (can be modified to avoid infinite loops)

61

DF-search; evaluation



- Completeness;
 - NO unless finite state space / depth
- Time complexity?
 - Assume maximum depth m

62

DF-search; evaluation



- Completeness;
 - NO unless finite state space / depth.
- Time complexity;
 - Tree search: worst case $O(b^m)$
 - Terrible if m is much larger than d (depth of optimal solution)
 - But if many solutions, may be much faster than BF-search

63

DF-search; evaluation



- Completeness;
 - NO unless finite state space / depth.
- Time complexity; $O(b^m)$
- Space complexity?

64

DF-search; evaluation



- Completeness;
 - NO unless finite state space / depth.
- Time complexity; $O(b^m)$
- Space complexity;
 - Tree search: linear space $O(bm)$; nodes expanded with no descendants in fringe can be removed from memory
 - Graph search: as Time, size of state space

65

DF-search; evaluation



- Completeness;
 - NO unless finite state space / depth.
- Time complexity; $O(b^m)$
- Space complexity;
 - Tree search: linear space $O(bm)$
- Optimality?

66

DF-search; evaluation



- Completeness;
 - NO unless finite state space / depth.
- Time complexity; $O(b^m)$
- Space complexity;
 - Tree search: linear space $O(bm)$
- Optimality; No

67

Depth-limited search



- DF-search with predetermined depth limit l .
 - i.e. nodes at depth l have no successors.
 - Problem knowledge can be used
- Solves the infinite-path problem.
- If $l < d$ then incompleteness results.
- not optimal.
- Time complexity: $O(b^l)$
- Space complexity: $O(bl)$

68

Depth-limited algorithm



```
function DEPTH-LIMITED-SEARCH(problem, limit) return a solution or failure/cutoff  
  return RECURSIVE-DLS(MAKE-NODE(problem.INITIAL-STATE),  
    problem, limit)
```

```
function RECURSIVE-DLS(node, problem, limit) return a solution or failure/cutoff  
  cutoff_occurred?  $\leftarrow$  false  
  if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)  
  else if DEPTH[node] == limit then return cutoff  
  else for each action in problem.ACTIONS(node.STATE) do  
    child  $\leftarrow$  Child-Node(problem, node, action)  
    result  $\leftarrow$  RECURSIVE-DLS(child, problem, limit)  
    if result == cutoff then cutoff_occurred?  $\leftarrow$  true  
    else if result  $\neq$  failure then return result  
  if cutoff_occurred? then return cutoff else return failure
```

Recursive implementation

69

Iterative deepening DF search



- Iterative deepening depth-first tree search
- A general strategy to find best depth limit l .
 - Gradually increasing the depth limit until a goal is found at depth d , the depth of the shallowest goal-node.
- Combines benefits of DF and BF search

70

Iterative deepening search



```
function ITERATIVE_DEEPENING_SEARCH(problem)  
  return a solution or failure
```

```
  for depth  $\leftarrow$  0 to  $\infty$  do  
    result  $\leftarrow$  DEPTH-LIMITED_SEARCH(problem,  
    depth)  
    if result  $\neq$  cutoff then return result
```

71

ID-search, example



- Limit=0



72

ID-search, example


- Limit=1



73

ID-search, example


- Limit=2



74

ID-search, example

- Limit=3



75

ID search, evaluation

- Complete?

76

ID search, evaluation



- Completeness:
 - YES
- Time complexity?

77

ID search, evaluation



- Completeness:
 - YES
 - Time complexity: $O(b^d)$
 - Algorithm seems costly due to repeated generation of certain states.
 - Node generation:
 - level d: once
 - level d-1: 2
 - level d-2: 3
 - ...
 - level 2: d-1
 - level 1: d
- Num. Comparison for b=10 and d=5 solution at far right
- $$N(IDS) = (d)b + (d-1)b^2 + \dots + (1)b^d$$
- $$N(BFS) = b + b^2 + \dots + b^d$$
- $$N(IDS) = 50 + 400 + 3000 + 20000 + 100000 = 123,450$$
- $$N(BFS) = 10 + 100 + 1000 + 10000 + 100000 = 111,110$$

78

ID search, evaluation



- Completeness:
 - YES
- Time complexity: $O(b^d)$
- Space complexity?

79

ID search, evaluation



- Completeness:
 - YES
- Time complexity: $O(b^d)$
- Space complexity: $O(bd)$
 - depth-first tree search
- Optimal?

80


ID search, evaluation



- Completeness:
 - YES
- Time complexity: $O(b^d)$
- Space complexity: $O(bd)$
- Optimality:
 - YES if step cost is 1.
 - Iterative analog to uniform-cost search?
- *ID is the preferred uninformed search method when there is a large search space and the depth of the solution is not known*

81


Bidirectional search



- Two simultaneous searches from start and goal.
 - Motivation: $b^{d/2} + b^{d/2} \ll b^d$
- Check whether the node belongs to the other fringe before expansion.
- Complete and optimal (for uniform step costs) if both searches are BF.
- Space complexity $O(b^{d/2})$ is the most significant weakness.

82

How to search backwards?



- Goal implicitly defined?
- The predecessor of each node should be efficiently computable.
 - When actions are easily reversible.

83

Summary of tree-search algorithms



Criterion	Breadth-First	Uniform-cost	Depth-First	Depth-limited	Iterative deepening	Bidirectional BF search
Complete?	YES	YES	NO	NO	YES	YES
Time	b^d	$b^{C*/e+1}$	b^m	b^l	b^d	$b^{d/2}$
Space	b^d	$b^{C*/e+1}$	bm	bl	bd	$b^{d/2}$
Optimal?	YES*	YES	NO	NO	YES*	YES*

* Optimal if identical step costs

84