

Outline

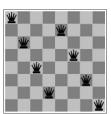


- Hill climbing
- Simulated annealing
- Local beam search
- Genetic algorithms (briefly)

Local search and optimization



- Previously: systematic exploration of search space.
 - solution to problem is path to goal
- Local search suitable for problems in which *path is irrelevant;* the goal state itself is the solution
- E.g 8-queens

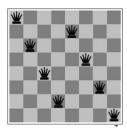


3

Local search and optimization



- Local search suitable for Optimization problems.
 - State space = set of "complete" configurations
 - Find best state according to some objective function h(s).
 - E.g., h(s) = # of conflicts
- E.g 8-queens using complete-state formulation
 - States: 8 queens on the board, one per column
 - Actions: Move a queen to a square in the same column



Local search and optimization



- Local search= keep a single current state and move to neighboring states to improve it
- Advantages:
 - Use very little memory
 - Often find reasonable solutions in large or infinite state spaces unsuitable for systematic algorithms
 - solve million-queens quickly

5

Hill-climbing search



- Keep a single current node and move to the best neighboring states to improve it
- "a loop that continuously moves in the direction of increasing value"
 - chooses randomly to break ties
 - It terminates when a peak is reached where no neighbor has a higher value
- Hill-climbing a.k.a. greedy local search, steepest ascent/descent

Hill-climbing search



function HILL-CLIMBING(*problem*) **return** a state that is a local maximum

 $\textit{current} \leftarrow \texttt{MAKE-NODE}(\texttt{problem.INITIAL-STATE}) \\ \textbf{loop do} \\$

 $neighbor \leftarrow$ a highest valued successor of current if $neighbor.VALUE \leq current.VALUE$ then return current.STATE $current \leftarrow neighbor$

7

Hill-climbing example



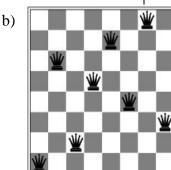
- 8-queens problem (complete-state formulation).
 - States: 8 queens on the board, one per column
 - Action: move a single queen to another square in the same column.
- Heuristic function h(n): the number of pairs of queens that are attacking each other (directly or indirectly).

Hill-climbing example



a)





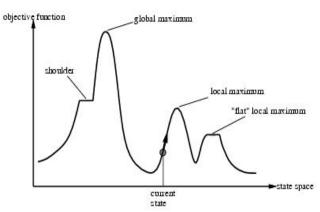
- a) shows a state of h=17 and the h-value for each possible successor.
- b) A local minimum in the 8-queens state space (h=1).

9

Drawbacks



state space landscape



Depending on initial state, can get stuck in local maxima, plateaux

Hill-climbing variations



- Stochastic hill-climbing
 - Random selection among the uphill moves.
 - The selection probability can vary with the steepness of the uphill move.
- First-choice hill-climbing
 - generating successors randomly until a better (than the current) one is found.
 - useful when a state has many successors
- Random-restart hill-climbing
 - Restart search from random initial state

11

Simulated annealing



- Escape local maxima by allowing "bad" moves.
 - Idea: but gradually decrease their frequency.
- T_r a "temperature" controlling the probability of downward steps

Simulated annealing



function SIMULATED-ANNEALING(*problem, schedule*) **return** a solution state

input: problem, a problem

schedule, a mapping from time to temperature

local variables: *T*, a "temperature" controlling the probability of downward steps

current ← MAKE-NODE(problem.INITIAL-STATE)

for $t \leftarrow 1$ to ∞ do

 $T \leftarrow schedule[t]$

if T = 0 then return current

 $next \leftarrow a$ randomly selected successor of *current*

 $\Delta E \leftarrow next.VALUE - current.VALUE$

if $\Delta E > 0$ then $current \leftarrow next$

else *current* \leftarrow *next* only with probability $e^{\Delta E/T}$

13

Simulated annealing



- Escape local maxima by allowing "bad" moves.
- One can prove: If T decreases slowly enough, then simulated annealing search will find a global optimum with probability approaching 1
- Commonly used: T ← cT with c constant close to, but smaller than, 1
- Applied for VLSI layout, airline scheduling, etc.

Local beam search



- Keep track of k states instead of one
 - Initially: k random states
 - Next: determine all successors of k states
 - If any of successors is goal → finished
 - Else select *k* best from successors and repeat.
- Major difference with random-restart search
 - Information is shared among *k* search threads.
- Can suffer from lack of diversity.
 - Stochastic variant: choose k successors randomly with probability proportional to state value difference.

15

Genetic algorithms



- Inspired by the process of biological evolution
- Start with k randomly generated states/individuals (population)
- States are scored by evaluation function (fitness function).
- At each step, the most fit states are selected (survival of the fittest) probabilistically: used as seeds for producing the next generation population -- the children or offspring, by means of operations such as crossover and mutation
- The process is repeated until sufficiently fit states are discovered -- the best state has a score exceeding a criterion
- They have been applied successfully to a variety of learning tasks and optimization problems

Genetic algorithm



function GENETIC_ALGORITHM(population, FITNESS-FN) **return** an individual

input: population, a set of individuals
FITNESS-FN, a function which determines the fitness of an individual

repeat

```
new_population ← empty set

loop for i from 1 to SIZE(population) do

x ← RANDOM_SELECTION(population, FITNESS_FN)
y ← RANDOM_SELECTION(population, FITNESS_FN)
child ← REPRODUCE(x,y)
if (small random probability) then child ← MUTATE(child)
add child to new_population
population ← new_population
until some individual is fit enough or enough time has elapsed
return the best individual in population
```

17

Genetic algorithms



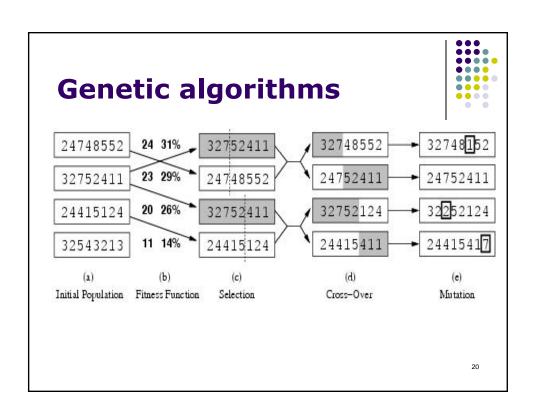
- In basic GAs, the fundamental representation of each state is a string over a finite alphabet (often a string of 0s and 1s), called a chromosome
- The mapping depends on the problem domain, and the designers.

Genetic algorithms



Genetic Operators

- Replication: a chromosome is merely reproduced
- Crossover: involves the mating of two parent chromosome to yield two new offspring by copying selected bits from each parent
- Mutation: creates a single descendant from a single parent by changing the value of a randomly chosen bit (from a 1 to 0 or vice versa)
- Other operators: specialized to the particular representation



Beyond Simple Environment



- Local Search in Continuous Spaces (Chapter 4.2)
- Searching with Nondeterministic Actions (Chapter 4.3)
- Searching with Partial Observations (Chapter 4.4)
- Online Search (Chapter 4.5)