

Games



Outline

- Games as search
- Optimal decisions in games
 - minimax decisions
 - α - β pruning
- Real-time decisions

Why study games?



- Examples: chess, checkers, Go, backgammon, bridge
- Why study games?
 - Fun; historically entertaining
 - Interesting subject of study because they are hard
 - Easy to represent and agents restricted to small number of actions
- Games are to AI as grand prix racing is to automobile design.

3

Games vs. Search



- Games are a form of *multi-agent environment*
 - What other agents do affect our success
- Competitive multi-agent environments give rise to *adversarial search*
- Solution is strategy (specifying a move for every possible opponent reply)

4

Types of Games



- Games of deterministic, perfect information: chess, checkers
- Stochastic games: backgammon
- Partially observable games: bridge, poker

5

Game setup



- Two players: MAX and MIN
- MAX moves first and they take turns until the game is over.
- Winner is awarded points (payoff), loser gets penalties.
- **Zero-sum game:** payoff values at the end of the game are equal and opposite (or the total payoff to all players is the same for every instance of the game)


6

Game setup

- Games as search:
 - Initial state: e.g. board configuration of chess
 - PLAYER(s): which player has the move in a state
 - Actions and transition model (or successor function): list of (move, state) pairs specifying legal moves.
 - Terminal test: Is the game finished (in *terminal states*)?
 - **Utility function** or **payoff function** $UTILITY(s,p)$: Gives numerical value for terminal states. E.g. win (+1), loss (-1), and draw (0) in tic-tac-toe
- The initial state and the legal moves define the **game tree**


7

Partial Game Tree for Tic-Tac-Toe



8

Two-Ply Game Tree



9

Optimal strategies




- Find the contingent *strategy* for MAX assuming an infallible MIN opponent.
- Assumption: Both players play optimally !!
- Given a game tree, the optimal strategy can be determined by using the **minimax value** of each node:

$$\text{MINIMAX-VALUE}(n) = \begin{cases} \text{UTILITY}(n) & \text{If } n \text{ is a terminal} \\ \max_{s \in \text{successors}(n)} \text{MINIMAX-VALUE}(s) & \text{If } n \text{ is a max node} \\ \min_{s \in \text{successors}(n)} \text{MINIMAX-VALUE}(s) & \text{If } n \text{ is a min node} \end{cases}$$

- Idea: choose move to position with highest minimax value
= best achievable payoff against best play


10

Two-Ply Game Tree



11

Two-Ply Game Tree



Minimax maximizes the worst-case outcome for max.

12

Minimax Algorithm



```
function MINIMAX-DECISION(state) returns an action
  inputs: state, current state in game
  return argmaxa in ACTIONS(state) MIN-VALUE(RESULT(state,a))
```

```
function MAX-VALUE(state) returns a utility value
  if TERMINAL-TEST(state) then return UTILITY(state)
  v ← - ∞
  for each a in ACTIONS(state) do
    v ← MAX(v, MIN-VALUE(RESULT(state,a)))
  return v
```

```
function MIN-VALUE(state) returns a utility value
  if TERMINAL-TEST(state) then return UTILITY(state)
  v ← ∞
  for each a in ACTIONS(state) do
    v ← MIN(v, MAX-VALUE(RESULT(state,a)))
  return v
```

13

Properties of minimax



- Complete? Yes (if tree is finite)
- Optimal? Yes, against an optimal opponent, otherwise?
- Time complexity? $O(b^m)$
- Space complexity? $O(bm)$ (depth-first exploration)
- For chess, $b \approx 35$, $m \approx 100$ for "reasonable" games
→ exact solution completely infeasible

14

What if MIN does not play optimally?



- Definition of optimal play for MAX assumes MIN plays optimally: maximizes worst-case outcome for MAX.
- But if MIN does not play optimally, MAX will do even better.
- There may be better strategies against suboptimal opponents, but they do worse against optimal opponents

15

Improving minimax search



- Number of states is exponential to the number of moves.
 - Solution: Do not examine every node
 - ==> Alpha-beta pruning
 - Remove branches that do not influence final decision
 - (Related algorithm: DFBnB)

16

Diversion: Depth-First Branch & Bound (DFBnB)




- Want to find the optimal solution among many solutions (or solving an optimization problem)
- DF search
- Keep track of best solution so far
- If $f(n) = g(n) + h(n) \geq \text{cost}(\text{best-soln})$
 - Then prune n

17


Diversion: DFBnB




Find the shortest path from V_0 to V_3



Diversion: DFBnB




Diversion: DFBnB



20

Pruning example




21

Alpha-Beta Example




Do DF-search until first leaf




22

Alpha-Beta Example (continued)




23

Alpha-Beta Example (continued)




24

Alpha-Beta Example (continued)




25

Alpha-Beta Example (continued)




26

Alpha-Beta Example (continued)



27

Alpha-Beta Example (continued)



28

Why is it called α - β ?



- Similar for MIN nodes
- α is the value of the best (i.e., highest-value) choice found so far at any choice point along the path for *max*
- β is the value of the best (i.e., lowest-value) choice found so far at any choice point along the path for *min*

29

The α - β algorithm



```
function ALPHA-BETA-SEARCH(state) returns an action
  inputs: state, current state in game
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$ 
  return the action in  $\text{ACTIONS}(\text{state})$  with value  $v$ 
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value
  if  $\text{TERMINAL-TEST}(\text{state})$  then return  $\text{UTILITY}(\text{state})$ 
   $v \leftarrow -\infty$ 
  for each  $a$  in  $\text{ACTIONS}(\text{state})$  do
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{Result}(s,a), \alpha, \beta))$ 
    if  $v \geq \beta$  then return  $v$ 
     $\alpha \leftarrow \text{MAX}(\alpha, v)$ 
  return  $v$ 
```

30

The α - β algorithm



```
function MIN-VALUE( $state, \alpha, \beta$ ) returns a utility value
  if TERMINAL-TEST( $state$ ) then return UTILITY( $state$ )
   $v \leftarrow +\infty$ 
  for each  $a$  in ACTIONS( $state$ ) do
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{Result}(s,a), \alpha, \beta))$ 
    if  $v \leq \alpha$  then return  $v$ 
   $\beta \leftarrow \text{MIN}(\beta, v)$ 
  return  $v$ 
```

31

Properties of Alpha-Beta Pruning



- Pruning does not affect final results
- The effectiveness of pruning is highly dependent on the order in which the successors are checked
 - Try to examine first the successors that are likely to be best
 - ordering heuristics
- With “perfect ordering,” time complexity is $O(b^{m/2})$
 - Branching factor of \sqrt{b} !!
 - Alpha-beta pruning can look twice as far as minimax in the same amount of time

32

Real-time Decisions



- Minimax and alpha-beta pruning require leaf-node evaluations.
- May be impractical within a reasonable amount of time.
- Suppose we have 100 secs, explore 10^6 nodes/sec
 - 10^8 nodes per move $\sim 35^5$
 - Alpha-beta pruning reaches depth 10

33

Real-time Decisions



Standard approach

- Cut off search earlier (replace TERMINAL-TEST by CUTOFF-TEST)
 - E.g., introduce a fixed depth limit (selected so that the amount of time will not exceed what the rules of the game allow), or iterative deepening
- Apply heuristic *evaluation function* EVAL (replacing utility function of alpha-beta)
- Change:
 - **if** TERMINAL-TEST(*state*) **then return** UTILITY(*state*)
 - into
 - **if** CUTOFF-TEST(*state*, *depth*) **then return** EVAL(*state*)

34

Heuristic evaluation function



- Idea: produce an estimate of the expected utility of the game from a given position.
- Performance depends strongly on quality of evaluation function
- Requirements:
 - EVAL should order terminal-nodes in the same way as UTILITY.
 - Computation may not take too long.
 - For non-terminal states the EVAL should be strongly correlated with the actual chance of winning.
- Exact values don't matter, only the order matters

35

Evaluation functions



- Most evaluation functions work by calculating various features of the state
 - Compute numerical contributions from each feature and combine them to find the total value
- For chess, *material value* for each piece: pawn=1, knight=bishop=3, rook=5, queen=9
- **Linear** weighted sum of **features**
 - $$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$
 - e.g., $w_1 = 9$ with $f_1(s) = (\text{number of white queens}) - (\text{number of black queens})$

36

Real-time Decisions



- Minimax with alpha-beta pruning
- An evaluation function
- More sophisticated cutoff test: e.g., wild swings in value in near future
- A large *transposition table* (closed list): repeated states problem

37

Real-time Decisions



- Extensively tuned evaluation function
- Pruning heuristics
- A large database of optimal opening and endgame moves: table lookup instead of search
 - Chess endgames with up to 7 pieces solved

38

Go

- Minimax not successful in Go



39

Monte Carlo tree search

- Find the most promising moves by playing out many random moves to the end, using techniques in Markov Decision Processes
- Doesn't necessarily require game-specific knowledge -> no need for evaluation function; simply implementing the game's mechanics is sufficient
- Does well in games with a high branching factor
- Used in Go, real-time video games
- AlphaGo: uses a Monte Carlo tree search guided by deep neural networks extensively trained from games played by human experts and games against itself
- AlphaGo Zero (2017) learns to play using Reinforcement Learning simply by playing games against itself, starting from completely random play

40

State-of-the-Art Deterministic games



- Checkers: Chinook ended 40-year-reign of human world champion Marion Tinsley in 1994. In July 2007, Chinook's developers announced that the program has been improved to the point where it cannot lose a game. Used a precomputed endgame database defining perfect play for all positions involving 8 or fewer pieces on the board, a total of 444 billion positions.
- Chess: Deep Blue defeated human world champion Garry Kasparov in a six-game match in 1997. Deep Blue searches 200 million positions per second, uses very sophisticated evaluation.
- Othello/Reversi: humans are no match for computers
- Go: AlphaGo beat a professional player without handicaps on a full-sized 19×19 board in 2015, beat a top Go player in 2016.

41

Types of Games



- We covered games of deterministic, perfect information: chess, checkers
- Stochastic Games: backgammon (Chapter 5.5)
- Partially observable games: bridge, poker (Chapter 5.6)

42